

Operations on NumPy Arrays

The learning objectives of this section are:

- Manipulate arrays
 - Reshape arrays
 - Stack arrays
- Perform operations on arrays
 - Perform basic mathematical operations
 - Apply built-in functions
 - Apply your own functions
 - Apply basic linear algebra operations

```
In [12]: import numpy as np
```

Example - 1 (Arithmetic Operations)

```
In [13]: array1 = np.array([10,20,30,40,50])  
array2 = np.arange(5)
```

```
In [14]: array1
```

```
Out[14]: array([10, 20, 30, 40, 50])
```

```
In [16]: array2
```

```
Out[16]: array([0, 1, 2, 3, 4])
```

```
In [17]: # Add array1 and array2.  
array3 = array1 + array2
```

```
In [18]: array3
```

```
Out[18]: array([10, 21, 32, 43, 54])
```

Example - 2

```
In [20]: array4 = np.array([1,2,3,4])
```

```
In [21]: array4 + array1
```

```
-----  
ValueError                                Traceback (most recent call last)  
<ipython-input-21-2811f702eb3f> in <module>  
----> 1 array4 + array1  
ValueError: operands could not be broadcast together with shapes (4,) (5,)
```

```
In [22]: print (array1.shape)  
(5,)
```

```
In [23]: print (array4.shape)  
(4,)
```

Example - 3

```
In [24]: array = np.linspace(1, 10, 5)  
array
```

```
Out[24]: array([ 1. ,  3.25,  5.5 ,  7.75, 10.  ])
```

```
In [25]: array*2
```

```
Out[25]: array([ 2. ,  6.5, 11. , 15.5, 20. ])
```

```
In [26]: array**2
```

```
Out[26]: array([ 1.    , 10.5625, 30.25 , 60.0625, 100.    ])
```

Stacking Arrays

`np.hstack()` and `np.vstack()`

Stacking is done using the `np.hstack()` and `np.vstack()` methods. For horizontal stacking, the number of rows should be the same, while for vertical stacking, the number of columns should be the same.

```
In [27]: # Note that np.hstack(a, b) throws an error - you need to pass the arrays as a list
a = np.array([1, 2, 3])
b = np.array([2, 3, 4])

np.hstack((a,b))
```

```
Out[27]: array([1, 2, 3, 2, 3, 4])
```

```
In [28]: np.vstack((a,b))
```

```
Out[28]: array([[1, 2, 3],
               [2, 3, 4]])
```

```
In [29]: np.arange(12)
```

```
Out[29]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
In [30]: np.arange(12).reshape(3,4)
```

```
Out[30]: array([[ 0,  1,  2,  3],
               [ 4,  5,  6,  7],
               [ 8,  9, 10, 11]])
```

```
In [31]: array1 = np.arange(12).reshape(3,4) #3x4
array2 = np.arange(20).reshape(5,4) #5x4
```

```
In [33]: print (array1, '\n', array2)
```

```
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
[[ 0  1  2  3]
 [ 4  5  6  7]
 [ 8  9 10 11]]
```

```
[12 13 14 15]
[16 17 18 19]]
```

```
In [34]: np.vstack((array1,array2))
```

```
Out[34]: array([[ 0,  1,  2,  3],
                [ 4,  5,  6,  7],
                [ 8,  9, 10, 11],
                [ 0,  1,  2,  3],
                [ 4,  5,  6,  7],
                [ 8,  9, 10, 11],
                [12, 13, 14, 15],
                [16, 17, 18, 19]])
```

Example - 4 (Numpy Built-in functions)

```
In [35]: array1
```

```
Out[35]: array([[ 0,  1,  2,  3],
                [ 4,  5,  6,  7],
                [ 8,  9, 10, 11]])
```

```
In [36]: np.power(array1, 3)
```

```
Out[36]: array([[ 0,  1,  8, 27],
                [64, 125, 216, 343],
                [512, 729, 1000, 1331]])
```

```
In [38]: np.arange(9).reshape(3,3)
```

```
Out[38]: array([[0, 1, 2],
                [3, 4, 5],
                [6, 7, 8]])
```

```
In [39]: x = np.array([-2,-1, 0, 1,2])
x
```

```
Out[39]: array([-2, -1,  0,  1,  2])
```

```
In [40]: abs(x)
```

```
Out[40]: array([2, 1, 0, 1, 2])
```

```
In [41]: np.absolute(x)
```

```
Out[41]: array([2, 1, 0, 1, 2])
```

Example - 5 (Trigonometric functions)

```
In [42]: np.pi
```

```
Out[42]: 3.141592653589793
```

```
In [43]: theta = np.linspace(0, np.pi, 5)
```

```
In [44]: theta
```

```
Out[44]: array([0.          , 0.78539816, 1.57079633, 2.35619449, 3.14159265])
```

```
In [45]: np.sin(theta)
```

```
Out[45]: array([0.00000000e+00, 7.07106781e-01, 1.00000000e+00, 7.07106781e-01,
 1.22464680e-16])
```

```
In [46]: np.cos(theta)
```

```
Out[46]: array([ 1.00000000e+00, 7.07106781e-01, 6.12323400e-17, -7.07106781e-01,
 -1.00000000e+00])
```

```
In [47]: np.tan(theta)
```

```
Out[47]: array([ 0.00000000e+00, 1.00000000e+00, 1.63312394e+16, -1.00000000e+00,
 -1.22464680e-16])
```

Example - 6 (Exponential and logarithmic functions)

```
In [48]: x = [1, 2, 3, 10]
x = np.array(x)
```

```
In [49]: np.exp(x) # e=2.718...
```

```
Out[49]: array([2.71828183e+00, 7.38905610e+00, 2.00855369e+01, 2.20264658e+04])
```

```
In [50]: # 2^1, 2^2, 2^3, 2^10  
np.exp2(x)
```

```
Out[50]: array([ 2.,  4.,  8., 1024.])
```

```
In [51]: np.power(x,3)
```

```
Out[51]: array([ 1,  8, 27, 1000])
```

```
In [52]: np.log(x)
```

```
Out[52]: array([0.          , 0.69314718, 1.09861229, 2.30258509])
```

```
In [53]: np.log2(x)
```

```
Out[53]: array([0.          , 1.          , 1.5849625 , 3.32192809])
```

```
In [54]: np.log10(x)
```

```
Out[54]: array([0.          , 0.30103   , 0.47712125, 1.          ])
```

```
In [ ]: np.log
```

Example - 7

```
In [57]: x = np.arange(5)  
x
```

```
Out[57]: array([0, 1, 2, 3, 4])
```

```
In [59]: y = x * 10  
y
```

```
Out[59]: array([ 0, 10, 20, 30, 40])
```

```
In [58]: y = np.empty(5)
```

```
y
```

```
Out[58]: array([ 1.00000000e+00,  7.07106781e-01,  6.12323400e-17, -7.07106781e-01,
                -1.00000000e+00])
```

```
In [61]: np.multiply(x, 12, out=y)
```

```
Out[61]: array([ 0, 12, 24, 36, 48])
```

```
In [62]: y
```

```
Out[62]: array([ 0, 12, 24, 36, 48])
```

```
In [63]: y = np.zeros(10)
y
```

```
Out[63]: array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0.])
```

```
In [65]: np.power(2, x, out=y[::2])
```

```
Out[65]: array([ 1.,  2.,  4.,  8., 16.])
```

```
In [66]: y
```

```
Out[66]: array([ 1.,  0.,  2.,  0.,  4.,  0.,  8.,  0., 16.,  0.])
```

Example - 8 (Aggregates)

```
In [67]: x = np.arange(1,6)
x
```

```
Out[67]: array([1, 2, 3, 4, 5])
```

```
In [69]: sum(x)
```

```
Out[69]: 15
```

```
In [68]: np.add.reduce(x)
```

```
Out[68]: 15
```

```
In [70]: np.add.accumulate(x)
```

```
Out[70]: array([ 1,  3,  6, 10, 15])
```

```
In [72]: np.multiply.accumulate(x)
```

```
Out[72]: array([ 1,  2,  6, 24, 120])
```

```
In [ ]:
```

Apply Basic Linear Algebra Operations

NumPy provides the `np.linalg` package to apply common linear algebra operations, such as:

- `np.linalg.inv` : Inverse of a matrix
- `np.linalg.det` : Determinant of a matrix
- `np.linalg.eig` : Eigenvalues and eigenvectors of a matrix

Also, you can multiple matrices using `np.dot(a, b)` .

```
In [73]: # np.linalg documentation
         help(np.linalg)
```

Help on package numpy.linalg in numpy:

NAME

numpy.linalg

DESCRIPTION

``numpy.linalg``
=====

The NumPy linear algebra functions rely on BLAS and LAPACK to provide efficient low level implementations of standard linear algebra algorithms. Those libraries may be provided by NumPy itself using C versions of a subset of their reference implementations but, when possible, highly optimized libraries that take advantage of specialized processor functionality are preferred. Examples

of such libraries are OpenBLAS, MKL (TM), and ATLAS. Because those libraries are multithreaded and processor dependent, environmental variables and external packages such as threadpoolctl may be needed to control the number of threads or specify the processor architecture.

- OpenBLAS: <https://www.openblas.net/>
- threadpoolctl: <https://github.com/joblib/threadpoolctl>

Please note that the most-used linear algebra functions in NumPy are present in the main `numpy` namespace rather than in `numpy.linalg`. There are: `dot`, `vdot`, `inner`, `outer`, `matmul`, `tensordot`, `einsum`, `einsum_path` and `kron`.

Functions present in `numpy.linalg` are listed below.

Matrix and vector products

`multi_dot`
`matrix_power`

Decompositions

`cholesky`
`qr`
`svd`

Matrix eigenvalues

`eig`
`eigh`
`eigvals`
`eigvalsh`

Norms and other numbers

`norm`
`cond`
`det`
`matrix_rank`
`slogdet`

Solving equations and inverting matrices

```
solve
tensorsolve
lstsq
inv
pinv
tensorinv
```

Exceptions

LinAlgError

PACKAGE CONTENTS

```
_umath_linalg
_lapack_lite
linalg
setup
tests (package)
```

CLASSES

```
builtins.Exception(builtins.BaseException)
    LinAlgError
```

```
class LinAlgError(builtins.Exception)
```

```
    Generic Python-exception-derived object raised by linalg functions.
```

```
    General purpose exception class, derived from Python's exception.Exception
    class, programmatically raised in linalg functions when a Linear
    Algebra-related condition would prevent further correct execution of the
    function.
```

```
    Parameters
```

```
    -----
```

```
    None
```

```
    Examples
```

```
    -----
```

```
>>> from numpy import linalg as LA
```

```
>>> LA.inv(np.zeros((2,2)))
```

```
Traceback (most recent call last):
```

```
  File "<stdin>", line 1, in <module>
```

```
File "...linalg.py", line 350,
    in inv return wrap(solve(a, identity(a.shape[0], dtype=a.dtype)))
File "...linalg.py", line 249,
    in solve
    raise LinAlgError('Singular matrix')
numpy.linalg.LinAlgError: Singular matrix

Method resolution order:
  LinAlgError
  builtins.Exception
  builtins.BaseException
  builtins.object

Data descriptors defined here:

__weakref__
    list of weak references to the object (if defined)
-----
Methods inherited from builtins.Exception:

__init__(self, /, *args, **kwargs)
    Initialize self.  See help(type(self)) for accurate signature.
-----
Static methods inherited from builtins.Exception:

__new__(*args, **kwargs) from builtins.type
    Create and return a new object.  See help(type) for accurate signature.
-----
Methods inherited from builtins.BaseException:

__delattr__(self, name, /)
    Implement delattr(self, name).

__getattr__(self, name, /)
    Return getattr(self, name).

__reduce__(...)
    Helper for pickle.

__repr__(self, /)
    Return repr(self).
```

```

__setattr__(self, name, value, /)
    Implement setattr(self, name, value).

__setstate__(...)

__str__(self, /)
    Return str(self).

with_traceback(...)
    Exception.with_traceback(tb) --
    set self.__traceback__ to tb and return self.

```

Data descriptors inherited from builtins.BaseException:

```

__cause__
    exception cause

__context__
    exception context

__dict__

__suppress_context__

__traceback__

args

```

FUNCTIONS

cholesky(a)
Cholesky decomposition.

Return the Cholesky decomposition, $L * L.H$, of the square matrix a , where L is lower-triangular and $.H$ is the conjugate transpose operator (which is the ordinary transpose if a is real-valued). a must be Hermitian (symmetric if real-valued) and positive-definite. Only L is actually returned.

Parameters

a : (... , M, M) array_like
 Hermitian (symmetric if all elements are real), positive-definite input matrix.

Returns

`L` : (... , M, M) array_like

Upper or lower-triangular Cholesky factor of ``a``. Returns a matrix object if ``a`` is a matrix object.

Raises

`LinAlgError`

If the decomposition fails, for example, if ``a`` is not positive-definite.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

The Cholesky decomposition is often used as a fast way of solving

.. $A \mathbf{x} = \mathbf{b}$

(when ``A`` is both Hermitian/symmetric and positive-definite).

First, we solve for \mathbf{y} in

.. $L \mathbf{y} = \mathbf{b}$,

and then for \mathbf{x} in

.. $L.H \mathbf{x} = \mathbf{y}$.

Examples

```
>>> A = np.array([[1,-2j],[2j,5]])
```

```
>>> A
```

```
array([[ 1.+0.j, -0.-2.j],
       [ 0.+2.j,  5.+0.j]])
```

```
>>> L = np.linalg.cholesky(A)
```

```
>>> L
```

```
array([[1.+0.j,  0.+0.j],
       [0.+2.j,  1.+0.j]])
```

```
>>> np.dot(L, L.T.conj()) # verify that L * L.H = A
```

```

array([[1.+0.j, 0.-2.j],
       [0.+2.j, 5.+0.j]])
>>> A = [[1,-2j],[2j,5]] # what happens if A is only array_like?
>>> np.linalg.cholesky(A) # an ndarray object is returned
array([[1.+0.j, 0.+0.j],
       [0.+2.j, 1.+0.j]])
>>> # But a matrix object is returned if A is a matrix object
>>> np.linalg.cholesky(np.matrix(A))
matrix([[ 1.+0.j,  0.+0.j],
        [ 0.+2.j,  1.+0.j]])

```

`cond(x, p=None)`

Compute the condition number of a matrix.

This function is capable of returning the condition number using one of seven different norms, depending on the value of `p` (see Parameters below).

Parameters

`x` : (... , M, N) array_like
The matrix whose condition number is sought.
`p` : {None, 1, -1, 2, -2, inf, -inf, 'fro'}, optional
Order of the norm:

```

=====
p      norm for matrices
=====
None    2-norm, computed directly using the ``SVD``
'fro'   Frobenius norm
inf     max(sum(abs(x), axis=1))
-inf    min(sum(abs(x), axis=1))
1       max(sum(abs(x), axis=0))
-1      min(sum(abs(x), axis=0))
2       2-norm (largest sing. value)
-2      smallest singular value
=====

```

inf means the `numpy.inf` object, and the Frobenius norm is the root-of-sum-of-squares norm.

Returns

`c` : {float, inf}
The condition number of the matrix. May be infinite.

See Also

`numpy.linalg.norm`

Notes

The condition number of ``x`` is defined as the norm of ``x`` times the norm of the inverse of ``x`` [1]_; the norm can be the usual L2-norm (root-of-sum-of-squares) or one of a number of other matrix norms.

References

.. [1] G. Strang, **Linear Algebra and Its Applications**, Orlando, FL, Academic Press, Inc., 1980, pg. 285.

Examples

```
>>> from numpy import linalg as LA
>>> a = np.array([[1, 0, -1], [0, 1, 0], [1, 0, 1]])
>>> a
array([[ 1,  0, -1],
       [ 0,  1,  0],
       [ 1,  0,  1]])
>>> LA.cond(a)
1.4142135623730951
>>> LA.cond(a, 'fro')
3.1622776601683795
>>> LA.cond(a, np.inf)
2.0
>>> LA.cond(a, -np.inf)
1.0
>>> LA.cond(a, 1)
2.0
>>> LA.cond(a, -1)
1.0
>>> LA.cond(a, 2)
1.4142135623730951
>>> LA.cond(a, -2)
0.70710678118654746 # may vary
>>> min(LA.svd(a, compute_uv=False))*min(LA.svd(LA.inv(a), compute_uv=False))
0.70710678118654746 # may vary
```

`det(a)`

Compute the determinant of an array.

Parameters

a : (... , M, M) array_like
Input array to compute determinants for.

Returns

det : (...) array_like
Determinant of `a`.

See Also

slogdet : Another way to represent the determinant, more suitable for large matrices where underflow/overflow may occur.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

The determinant is computed via LU factorization using the LAPACK routine ``z/dgetrf``.

Examples

The determinant of a 2-D array `[[a, b], [c, d]]` is `ad - bc`:

```
>>> a = np.array([[1, 2], [3, 4]])
>>> np.linalg.det(a)
-2.0 # may vary
```

Computing determinants for a stack of matrices:

```
>>> a = np.array([ [1, 2], [3, 4]], [1, 2], [2, 1], [1, 3], [3, 1] ])
>>> a.shape
(3, 2, 2)
>>> np.linalg.det(a)
array([-2., -3., -8.])
```

eig(a)

Compute the eigenvalues and right eigenvectors of a square array.

Parameters

a : (... , M, M) array

Matrices for which the eigenvalues and right eigenvectors will be computed

Returns

w : (... , M) array

The eigenvalues, each repeated according to its multiplicity. The eigenvalues are not necessarily ordered. The resulting array will be of complex type, unless the imaginary part is zero in which case it will be cast to a real type. When `a` is real the resulting eigenvalues will be real (0 imaginary part) or occur in conjugate pairs

v : (... , M, M) array

The normalized (unit "length") eigenvectors, such that the column ``v[:,i]`` is the eigenvector corresponding to the eigenvalue ``w[i]``.

Raises

LinAlgError

If the eigenvalue computation does not converge.

See Also

eigvals : eigenvalues of a non-symmetric array.

eigh : eigenvalues and eigenvectors of a real symmetric or complex Hermitian (conjugate symmetric) array.

eigvalsh : eigenvalues of a real symmetric or complex Hermitian (conjugate symmetric) array.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the `numpy.linalg` documentation for details.

This is implemented using the ``_geev`` LAPACK routines which compute the eigenvalues and eigenvectors of general square arrays.

The number w is an eigenvalue of a if there exists a vector v such that $\text{dot}(a,v) = w * v$. Thus, the arrays a , w , and v satisfy the equations $\text{dot}(a[:,i], v[:,i]) = w[i] * v[:,i]$ for $i \in \{0, \dots, M-1\}$.

The array v of eigenvectors may not be of maximum rank, that is, some of the columns may be linearly dependent, although round-off error may obscure that fact. If the eigenvalues are all different, then theoretically the eigenvectors are linearly independent. Likewise, the (complex-valued) matrix of eigenvectors v is unitary if the matrix a is normal, i.e., if $\text{dot}(a, a.H) = \text{dot}(a.H, a)$, where $a.H$ denotes the conjugate transpose of a .

Finally, it is emphasized that v consists of the *right* (as in right-hand side) eigenvectors of a . A vector y satisfying $\text{dot}(y.T, a) = z * y.T$ for some number z is called a *left* eigenvector of a , and, in general, the left and right eigenvectors of a matrix are not necessarily the (perhaps conjugate) transposes of each other.

References

G. Strang, *Linear Algebra and Its Applications*, 2nd Ed., Orlando, FL, Academic Press, Inc., 1980, Various pp.

Examples

```
>>> from numpy import linalg as LA
```

(Almost) trivial example with real e-values and e-vectors.

```
>>> w, v = LA.eig(np.diag((1, 2, 3)))
>>> w; v
array([1., 2., 3.])
array([[1., 0., 0.],
       [0., 1., 0.],
       [0., 0., 1.]])
```

Real matrix possessing complex e-values and e-vectors; note that the e-values are complex conjugates of each other.

```
>>> w, v = LA.eig(np.array([[1, -1], [1, 1]]))
```

```
>>> w; v
array([1.+1.j, 1.-1.j])
array([[0.70710678+0.j, 0.70710678-0.j],
       [0. -0.70710678j, 0. +0.70710678j]])
```

Complex-valued matrix with real e-values (but complex-valued e-vectors);
note that ``a.conj().T == a``, i.e., ``a`` is Hermitian.

```
>>> a = np.array([[1, 1j], [-1j, 1]])
>>> w, v = LA.eig(a)
>>> w; v
array([2.+0.j, 0.+0.j])
array([[ 0. +0.70710678j, 0.70710678+0.j], # may vary
       [ 0.70710678+0.j, -0. +0.70710678j]])
```

Be careful about round-off error!

```
>>> a = np.array([[1 + 1e-9, 0], [0, 1 - 1e-9]])
>>> # Theor. e-values are 1 +/- 1e-9
>>> w, v = LA.eig(a)
>>> w; v
array([1., 1.])
array([[1., 0.],
       [0., 1.]])
```

`eigh(a, UPL0='L')`

Return the eigenvalues and eigenvectors of a complex Hermitian
(conjugate symmetric) or a real symmetric matrix.

Returns two objects, a 1-D array containing the eigenvalues of ``a``, and
a 2-D square array or matrix (depending on the input type) of the
corresponding eigenvectors (in columns).

Parameters

`a` : (... , M, M) array

Hermitian or real symmetric matrices whose eigenvalues and
eigenvectors are to be computed.

`UPL0` : {'L', 'U'}, optional

Specifies whether the calculation is done with the lower triangular
part of ``a`` ('L', default) or the upper triangular part ('U').

Irrespective of this value only the real parts of the diagonal will
be considered in the computation to preserve the notion of a Hermitian
matrix. It therefore follows that the imaginary part of the diagonal
will always be treated as zero.

Returns

w : (... , M) ndarray

The eigenvalues in ascending order, each repeated according to its multiplicity.

v : {... , M, M) ndarray, (... , M, M) matrix}

The column ``v[:, i]`` is the normalized eigenvector corresponding to the eigenvalue ``w[i]``. Will return a matrix object if ``a`` is a matrix object.

Raises

LinAlgError

If the eigenvalue computation does not converge.

See Also

eigvalsh : eigenvalues of real symmetric or complex Hermitian (conjugate symmetric) arrays.

eig : eigenvalues and right eigenvectors for non-symmetric arrays.

eigvals : eigenvalues of non-symmetric arrays.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

The eigenvalues/eigenvectors are computed using LAPACK routines ``_syevd``, ``_heevd``.

The eigenvalues of real symmetric or complex Hermitian matrices are always real. [1]_ The array ``v`` of (column) eigenvectors is unitary and ``a``, ``w``, and ``v`` satisfy the equations
``dot(a, v[:, i]) = w[i] * v[:, i]``.

References

.. [1] G. Strang, **Linear Algebra and Its Applications**, 2nd Ed., Orlando, FL, Academic Press, Inc., 1980, pg. 222.

Examples

```

-----
>>> from numpy import linalg as LA
>>> a = np.array([[1, -2j], [2j, 5]])
>>> a
array([[ 1.+0.j, -0.-2.j],
       [ 0.+2.j,  5.+0.j]])
>>> w, v = LA.eigh(a)
>>> w; v
array([0.17157288, 5.82842712])
array([[ -0.92387953+0.j, -0.38268343+0.j], # may vary
       [ 0.          +0.38268343j,  0.          -0.92387953j]])

>>> np.dot(a, v[:, 0]) - w[0] * v[:, 0] # verify 1st e-val/vec pair
array([5.55111512e-17+0.00000000e+00j, 0.00000000e+00+1.2490009e-16j])
>>> np.dot(a, v[:, 1]) - w[1] * v[:, 1] # verify 2nd e-val/vec pair
array([0.+0.j, 0.+0.j])

>>> A = np.matrix(a) # what happens if input is a matrix object
>>> A
matrix([[ 1.+0.j, -0.-2.j],
        [ 0.+2.j,  5.+0.j]])
>>> w, v = LA.eigh(A)
>>> w; v
array([0.17157288, 5.82842712])
matrix([[ -0.92387953+0.j, -0.38268343+0.j], # may vary
        [ 0.          +0.38268343j,  0.          -0.92387953j]])

>>> # demonstrate the treatment of the imaginary part of the diagonal
>>> a = np.array([[5+2j, 9-2j], [0+2j, 2-1j]])
>>> a
array([[5.+2.j, 9.-2.j],
       [0.+2.j, 2.-1.j]])
>>> # with UPL0='L' this is numerically equivalent to using LA.eig() with:
>>> b = np.array([[5.+0.j, 0.-2.j], [0.+2.j, 2.-0.j]])
>>> b
array([[5.+0.j, 0.-2.j],
       [0.+2.j, 2.+0.j]])
>>> wa, va = LA.eigh(a)
>>> wb, vb = LA.eig(b)
>>> wa; wb
array([1., 6.])
array([6.+0.j, 1.+0.j])
>>> va; vb
array([[ -0.4472136 +0.j, -0.89442719+0.j], # may vary
       [ 0.          +0.89442719j,  0.          -0.4472136j ]])

```

```
array([[ 0.89442719+0.j, -0.         +0.4472136j],
       [-0.         +0.4472136j,  0.89442719+0.j   ]])
```

`eigvals(a)`

Compute the eigenvalues of a general matrix.

Main difference between ``eigvals`` and ``eig``: the eigenvectors aren't returned.

Parameters

`a` : (... , M, M) array_like

A complex- or real-valued matrix whose eigenvalues will be computed.

Returns

`w` : (... , M,) ndarray

The eigenvalues, each repeated according to its multiplicity.
They are not necessarily ordered, nor are they necessarily
real for real matrices.

Raises

`LinAlgError`

If the eigenvalue computation does not converge.

See Also

`eig` : eigenvalues and right eigenvectors of general arrays

`eigvalsh` : eigenvalues of real symmetric or complex Hermitian
(conjugate symmetric) arrays.

`eigh` : eigenvalues and eigenvectors of real symmetric or complex
Hermitian (conjugate symmetric) arrays.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

This is implemented using the ``_geev`` LAPACK routines which compute the eigenvalues and eigenvectors of general square arrays.

Examples

Illustration, using the fact that the eigenvalues of a diagonal matrix are its diagonal elements, that multiplying a matrix on the left by an orthogonal matrix, `Q`, and on the right by `Q.T` (the transpose of `Q`), preserves the eigenvalues of the "middle" matrix. In other words, if `Q` is orthogonal, then ``Q * A * Q.T`` has the same eigenvalues as ``A``:

```
>>> from numpy import linalg as LA
>>> x = np.random.random()
>>> Q = np.array([[np.cos(x), -np.sin(x)], [np.sin(x), np.cos(x)]])
>>> LA.norm(Q[0, :]), LA.norm(Q[1, :]), np.dot(Q[0, :], Q[1, :])
(1.0, 1.0, 0.0)
```

Now multiply a diagonal matrix by ``Q`` on one side and by ``Q.T`` on the other:

```
>>> D = np.diag((-1,1))
>>> LA.eigvals(D)
array([-1.,  1.])
>>> A = np.dot(Q, D)
>>> A = np.dot(A, Q.T)
>>> LA.eigvals(A)
array([ 1., -1.]) # random
```

`eigvalsh(a, UPL0='L')`

Compute the eigenvalues of a complex Hermitian or real symmetric matrix.

Main difference from `eigh`: the eigenvectors are not computed.

Parameters

`a` : (... , M, M) array_like

A complex- or real-valued matrix whose eigenvalues are to be computed.

`UPL0` : {'L', 'U'}, optional

Specifies whether the calculation is done with the lower triangular part of `a` ('L', default) or the upper triangular part ('U').

Irrespective of this value only the real parts of the diagonal will be considered in the computation to preserve the notion of a Hermitian matrix. It therefore follows that the imaginary part of the diagonal will always be treated as zero.

Returns

w : (... , M,) ndarray
The eigenvalues in ascending order, each repeated according to its multiplicity.

Raises

LinAlgError

If the eigenvalue computation does not converge.

See Also

eigh : eigenvalues and eigenvectors of real symmetric or complex Hermitian (conjugate symmetric) arrays.

eigvals : eigenvalues of general real or complex arrays.

eig : eigenvalues and right eigenvectors of general real or complex arrays.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

The eigenvalues are computed using LAPACK routines ```_syevd```, ```_heevd```.

Examples

```
>>> from numpy import linalg as LA
```

```
>>> a = np.array([[1, -2j], [2j, 5]])
```

```
>>> LA.eigvalsh(a)
```

```
array([ 0.17157288,  5.82842712]) # may vary
```

```
>>> # demonstrate the treatment of the imaginary part of the diagonal
```

```
>>> a = np.array([[5+2j, 9-2j], [0+2j, 2-1j]])
```

```
>>> a
```

```
array([[5.+2.j, 9.-2.j],  
       [0.+2.j, 2.-1.j]])
```

```
>>> # with UPL0='L' this is numerically equivalent to using LA.eigvals()
```

```
>>> # with:
```

```
>>> b = np.array([[5.+0.j, 0.-2.j], [0.+2.j, 2.-0.j]])
```

```
>>> b
```

```
array([[5.+0.j, 0.-2.j],  
       [0.+2.j, 2.+0.j]])
```



```
>>> wa = LA.eigvalsh(a)
>>> wb = LA.eigvals(b)
>>> wa; wb
array([1., 6.])
array([6.+0.j, 1.+0.j])
```

`inv(a)`

Compute the (multiplicative) inverse of a matrix.

Given a square matrix ``a``, return the matrix ``ainv`` satisfying
``dot(a, ainv) = dot(ainv, a) = eye(a.shape[0])``.

Parameters

`a : (..., M, M) array_like`
Matrix to be inverted.

Returns

`ainv : (..., M, M) ndarray or matrix`
(Multiplicative) inverse of the matrix ``a``.

Raises

`LinAlgError`

If ``a`` is not square or inversion fails.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

Examples

```
>>> from numpy.linalg import inv
>>> a = np.array([[1., 2.], [3., 4.]])
>>> ainv = inv(a)
>>> np.allclose(np.dot(a, ainv), np.eye(2))
True
>>> np.allclose(np.dot(ainv, a), np.eye(2))
True
```

If `a` is a matrix object, then the return value is a matrix as well:

```
>>> ainv = inv(np.matrix(a))
>>> ainv
matrix([[ -2. ,  1. ],
        [ 1.5, -0.5]])
```

Inverses of several matrices can be computed at once:

```
>>> a = np.array([[[1., 2.], [3., 4.]], [[1, 3], [3, 5]]])
>>> inv(a)
array([[[ -2. ,  1. ],
        [ 1.5 , -0.5 ]],
       [[-1.25,  0.75],
        [ 0.75, -0.25]]])
```

`lstsq(a, b, rcond='warn')`

Return the least-squares solution to a linear matrix equation.

Solves the equation $a x = b$ by computing a vector x that minimizes the squared Euclidean 2-norm $\|b - a x\|_2^2$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the "exact" solution of the equation.

Parameters

`a` : (M, N) array_like

"Coefficient" matrix.

`b` : {(M,), (M, K)} array_like

Ordinate or "dependent variable" values. If b is two-dimensional, the least-squares solution is calculated for each of the K columns of b .

`rcond` : float, optional

Cut-off ratio for small singular values of a .

For the purposes of rank determination, singular values are treated as zero if they are smaller than `rcond` times the largest singular value of a .

.. versionchanged:: 1.14.0

If not set, a FutureWarning is given. The previous default of `-1` will use the machine precision as `rcond` parameter, the new default will use the machine precision times `max(M, N)`.

To silence the warning and use the new default, use ``rcond=None``,
to keep using the old behavior, use ``rcond=-1``.

Returns

```
-----
x : {(N,), (N, K)} ndarray
    Least-squares solution. If `b` is two-dimensional,
    the solutions are in the `K` columns of `x`.
residuals : {(1,), (K,), (0,)} ndarray
    Sums of residuals; squared Euclidean 2-norm for each column in
    ``b - a*x``.
    If the rank of `a` is < N or M <= N, this is an empty array.
    If `b` is 1-dimensional, this is a (1,) shape array.
    Otherwise the shape is (K,).
rank : int
    Rank of matrix `a`.
s : (min(M, N),) ndarray
    Singular values of `a`.
```

Raises

```
-----
LinAlgError
    If computation does not converge.
```

Notes

```
-----
If `b` is a matrix, then all array results are returned as matrices.
```

Examples

```
-----
Fit a line, ``y = mx + c``, through some noisy data-points:
```

```
>>> x = np.array([0, 1, 2, 3])
>>> y = np.array([-1, 0.2, 0.9, 2.1])
```

By examining the coefficients, we see that the line should have a
gradient of roughly 1 and cut the y-axis at, more or less, -1.

We can rewrite the line equation as ``y = Ap``, where ``A = [[x 1]]``
and ``p = [[m], [c]]``. Now use `lstsq` to solve for `p`:

```
>>> A = np.vstack([x, np.ones(len(x))]).T
>>> A
array([[ 0.,  1.],
       [ 1.,  1.]
```

```

        [ 2.,  1.],
        [ 3.,  1.]])

>>> m, c = np.linalg.lstsq(A, y, rcond=None)[0]
>>> m, c
(1.0 -0.95) # may vary

```

Plot the data along with the fitted line:

```

>>> import matplotlib.pyplot as plt
>>> _ = plt.plot(x, y, 'o', label='Original data', markersize=10)
>>> _ = plt.plot(x, m*x + c, 'r', label='Fitted line')
>>> _ = plt.legend()
>>> plt.show()

```

`matrix_power(a, n)`

Raise a square matrix to the (integer) power `n`.

For positive integers `n`, the power is computed by repeated matrix squarings and matrix multiplications. If `n == 0`, the identity matrix of the same shape as `M` is returned. If `n < 0`, the inverse is computed and then raised to the `abs(n)`.

.. note:: Stacks of object matrices are not currently supported.

Parameters

`a` : (... , M, M) array_like
Matrix to be "powered".

`n` : int
The exponent can be any integer or long integer, positive, negative, or zero.

Returns

`a**n` : (... , M, M) ndarray or matrix object
The return value is the same shape and type as `M`;
if the exponent is positive or zero then the type of the elements is the same as those of `M`. If the exponent is negative the elements are floating-point.

Raises

`LinAlgError`

For matrices that are not square or that (for negative powers) cannot

be inverted numerically.

Examples

```
>>> from numpy.linalg import matrix_power
>>> i = np.array([[0, 1], [-1, 0]]) # matrix equiv. of the imaginary unit
>>> matrix_power(i, 3) # should = -i
array([[ 0, -1],
       [ 1,  0]])
>>> matrix_power(i, 0)
array([[1, 0],
       [0, 1]])
>>> matrix_power(i, -3) # should = 1/(-i) = i, but w/ f.p. elements
array([[ 0.,  1.],
       [-1.,  0.]])
```

Somewhat more sophisticated example

```
>>> q = np.zeros((4, 4))
>>> q[0:2, 0:2] = -i
>>> q[2:4, 2:4] = i
>>> q # one of the three quaternion units not equal to 1
array([[ 0., -1.,  0.,  0.],
       [ 1.,  0.,  0.,  0.],
       [ 0.,  0.,  0.,  1.],
       [ 0.,  0., -1.,  0.]])
>>> matrix_power(q, 2) # = -np.eye(4)
array([[ -1.,  0.,  0.,  0.],
       [ 0., -1.,  0.,  0.],
       [ 0.,  0., -1.,  0.],
       [ 0.,  0.,  0., -1.]])
```

`matrix_rank(M, tol=None, hermitian=False)`
Return matrix rank of array using SVD method

Rank of the array is the number of singular values of the array that are greater than `tol`.

.. versionchanged:: 1.14
Can now operate on stacks of matrices

Parameters

`M : {(M,)}, (... , M, N)} array_like`
Input vector or stack of matrices.

```

tol : (...) array_like, float, optional
    Threshold below which SVD values are considered zero. If `tol` is
    None, and ``S`` is an array with singular values for `M`, and
    ``eps`` is the epsilon value for datatype of ``S``, then `tol` is
    set to ``S.max() * max(M.shape) * eps``.

    .. versionchanged:: 1.14
        Broadcasted against the stack of matrices
hermitian : bool, optional
    If True, `M` is assumed to be Hermitian (symmetric if real-valued),
    enabling a more efficient method for finding singular values.
    Defaults to False.

    .. versionadded:: 1.14

```

Returns

```

-----
rank : (...) array_like
    Rank of M.

```

Notes

```

-----

```

The default threshold to detect rank deficiency is a test on the magnitude of the singular values of `M`. By default, we identify singular values less than ``S.max() * max(M.shape) * eps`` as indicating rank deficiency (with the symbols defined above). This is the algorithm MATLAB uses [1]. It also appears in *Numerical recipes* in the discussion of SVD solutions for linear least squares [2].

This default threshold is designed to detect rank deficiency accounting for the numerical errors of the SVD computation. Imagine that there is a column in `M` that is an exact (in floating point) linear combination of other columns in `M`. Computing the SVD on `M` will not produce a singular value exactly equal to 0 in general: any difference of the smallest SVD value from 0 will be caused by numerical imprecision in the calculation of the SVD. Our threshold for small SVD values takes this numerical imprecision into account, and the default threshold will detect such numerical rank deficiency. The threshold may declare a matrix `M` rank deficient even if the linear combination of some columns of `M` is not exactly equal to another column of `M` but only numerically very close to another column of `M`.

We chose our default threshold because it is in wide use. Other thresholds are possible. For example, elsewhere in the 2007 edition of *Numerical recipes* there is an alternative threshold of ``S.max() *

`np.finfo(M.dtype).eps / 2. * np.sqrt(m + n + 1.)```. The authors describe this threshold as being based on "expected roundoff error" (p 71).

The thresholds above deal with floating point roundoff error in the calculation of the SVD. However, you may have more information about the sources of error in `M` that would make you consider other tolerance values to detect *effective* rank deficiency. The most useful measure of the tolerance depends on the operations you intend to use on your matrix. For example, if your data come from uncertain measurements with uncertainties greater than floating point epsilon, choosing a tolerance near that uncertainty may be preferable. The tolerance may be absolute if the uncertainties are absolute rather than relative.

References

-
- .. [1] MATLAB reference documentation, "Rank"
<https://www.mathworks.com/help/techdoc/ref/rank.html>
 - .. [2] W. H. Press, S. A. Teukolsky, W. T. Vetterling and B. P. Flannery, "Numerical Recipes (3rd edition)", Cambridge University Press, 2007, page 795.

Examples

```
>>> from numpy.linalg import matrix_rank
>>> matrix_rank(np.eye(4)) # Full rank matrix
4
>>> I=np.eye(4); I[-1,-1] = 0. # rank deficient matrix
>>> matrix_rank(I)
3
>>> matrix_rank(np.ones((4,))) # 1 dimension - rank 1 unless all 0
1
>>> matrix_rank(np.zeros((4,)))
0
```

`multi_dot(arrays)`

Compute the dot product of two or more arrays in a single function call, while automatically selecting the fastest evaluation order.

`multi_dot` chains `numpy.dot` and uses optimal parenthesization of the matrices [1]_ [2]_. Depending on the shapes of the matrices, this can speed up the multiplication a lot.

If the first argument is 1-D it is treated as a row vector.
If the last argument is 1-D it is treated as a column vector.
The other arguments must be 2-D.

Think of ``multi_dot`` as::

```
def multi_dot(arrays): return functools.reduce(np.dot, arrays)
```

Parameters

`arrays` : sequence of array_like

If the first argument is 1-D it is treated as row vector.

If the last argument is 1-D it is treated as column vector.

The other arguments must be 2-D.

Returns

`output` : ndarray

Returns the dot product of the supplied arrays.

See Also

`dot` : dot multiplication with two arguments.

References

.. [1] Cormen, "Introduction to Algorithms", Chapter 15.2, p. 370-378

.. [2] https://en.wikipedia.org/wiki/Matrix_chain_multiplication

Examples

``multi_dot`` allows you to write::

```
>>> from numpy.linalg import multi_dot
>>> # Prepare some data
>>> A = np.random.random((10000, 100))
>>> B = np.random.random((100, 1000))
>>> C = np.random.random((1000, 5))
>>> D = np.random.random((5, 333))
>>> # the actual dot multiplication
>>> _ = multi_dot([A, B, C, D])
```

instead of::

```
>>> _ = np.dot(np.dot(np.dot(A, B), C), D)
>>> # or
```



```
>>> _ = A.dot(B).dot(C).dot(D)
```

Notes

The cost for a matrix multiplication can be calculated with the following function::

```
def cost(A, B):  
    return A.shape[0] * A.shape[1] * B.shape[1]
```

Assume we have three matrices

:math:`A_{\{10 \times 100\}}, B_{\{100 \times 5\}}, C_{\{5 \times 50\}}`.

The costs for the two different parenthesizations are as follows::

```
cost((AB)C) = 10*100*5 + 10*5*50 = 5000 + 2500 = 7500  
cost(A(BC)) = 10*100*50 + 100*5*50 = 50000 + 25000 = 75000
```

`norm(x, ord=None, axis=None, keepdims=False)`
Matrix or vector norm.

This function is able to return one of eight different matrix norms, or one of an infinite number of vector norms (described below), depending on the value of the ``ord`` parameter.

Parameters

`x` : array_like

Input array. If ``axis`` is None, ``x`` must be 1-D or 2-D, unless ``ord`` is None. If both ``axis`` and ``ord`` are None, the 2-norm of ``x.ravel`` will be returned.

`ord` : {non-zero int, inf, -inf, 'fro', 'nuc'}, optional

Order of the norm (see table under ``Notes``). inf means numpy's ``inf`` object. The default is None.

`axis` : {None, int, 2-tuple of ints}, optional.

If ``axis`` is an integer, it specifies the axis of ``x`` along which to compute the vector norms. If ``axis`` is a 2-tuple, it specifies the axes that hold 2-D matrices, and the matrix norms of these matrices are computed. If ``axis`` is None then either a vector norm (when ``x`` is 1-D) or a matrix norm (when ``x`` is 2-D) is returned. The default is None.

.. versionadded:: 1.8.0

`keepdims` : bool, optional

If this is set to True, the axes which are normed over are left in the result as dimensions with size one. With this option the result will broadcast correctly against the original `x`.

.. versionadded:: 1.10.0

Returns

n : float or ndarray
Norm of the matrix or vector(s).

Notes

For values of ``ord <= 0``, the result is, strictly speaking, not a mathematical 'norm', but it may still be useful for various numerical purposes.

The following norms can be calculated:

ord	norm for matrices	norm for vectors
None	Frobenius norm	2-norm
'fro'	Frobenius norm	--
'nuc'	nuclear norm	--
inf	max(sum(abs(x), axis=1))	max(abs(x))
-inf	min(sum(abs(x), axis=1))	min(abs(x))
0	--	sum(x != 0)
1	max(sum(abs(x), axis=0))	as below
-1	min(sum(abs(x), axis=0))	as below
2	2-norm (largest sing. value)	as below
-2	smallest singular value	as below
other	--	sum(abs(x)**ord)**(1./ord)

The Frobenius norm is given by [1]_:

$$\|A\|_F = \left(\sum_{i,j} \text{abs}(a_{i,j})^2 \right)^{1/2}$$

The nuclear norm is the sum of the singular values.

References

.. [1] G. H. Golub and C. F. Van Loan, *Matrix Computations*,
Baltimore, MD, Johns Hopkins University Press, 1985, pg. 15

Examples

```
-----
>>> from numpy import linalg as LA
>>> a = np.arange(9) - 4
>>> a
array([-4, -3, -2, ..., 2, 3, 4])
>>> b = a.reshape((3, 3))
>>> b
array([[ -4,  -3,  -2],
       [ -1,   0,   1],
       [ 2,   3,   4]])

>>> LA.norm(a)
7.745966692414834
>>> LA.norm(b)
7.745966692414834
>>> LA.norm(b, 'fro')
7.745966692414834
>>> LA.norm(a, np.inf)
4.0
>>> LA.norm(b, np.inf)
9.0
>>> LA.norm(a, -np.inf)
0.0
>>> LA.norm(b, -np.inf)
2.0

>>> LA.norm(a, 1)
20.0
>>> LA.norm(b, 1)
7.0
>>> LA.norm(a, -1)
-4.6566128774142013e-010
>>> LA.norm(b, -1)
6.0
>>> LA.norm(a, 2)
7.745966692414834
>>> LA.norm(b, 2)
7.3484692283495345

>>> LA.norm(a, -2)
0.0
>>> LA.norm(b, -2)
1.8570331885190563e-016 # may vary
```

```
>>> LA.norm(a, 3)
5.8480354764257312 # may vary
>>> LA.norm(a, -3)
0.0
```

Using the `axis` argument to compute vector norms:

```
>>> c = np.array([[ 1, 2, 3],
...               [-1, 1, 4]])
>>> LA.norm(c, axis=0)
array([ 1.41421356,  2.23606798,  5.          ])
>>> LA.norm(c, axis=1)
array([ 3.74165739,  4.24264069])
>>> LA.norm(c, ord=1, axis=1)
array([ 6.,  6.])
```

Using the `axis` argument to compute matrix norms:

```
>>> m = np.arange(8).reshape(2,2,2)
>>> LA.norm(m, axis=(1,2))
array([ 3.74165739, 11.22497216])
>>> LA.norm(m[0, :, :], LA.norm(m[1, :, :])
(3.7416573867739413, 11.224972160321824)
```

`pinv(a, rcond=1e-15, hermitian=False)`
 Compute the (Moore-Penrose) pseudo-inverse of a matrix.

Calculate the generalized inverse of a matrix using its singular-value decomposition (SVD) and including all *large* singular values.

.. versionchanged:: 1.14
 Can now operate on stacks of matrices

Parameters

`a` : (...) array_like
 Matrix or stack of matrices to be pseudo-inverted.
`rcond` : (...) array_like of float
 Cutoff for small singular values.
 Singular values less than or equal to
 ``rcond * largest_singular_value`` are set to zero.
 Broadcasts against the stack of matrices.
`hermitian` : bool, optional
 If True, `a` is assumed to be Hermitian (symmetric if real-valued),

enabling a more efficient method for finding singular values.
Defaults to False.

.. versionadded:: 1.17.0

Returns

B : (... , N, M) ndarray

The pseudo-inverse of `a`. If `a` is a `matrix` instance, then so is `B`.

Raises

LinAlgError

If the SVD computation does not converge.

Notes

The pseudo-inverse of a matrix A , denoted A^+ , is defined as: "the matrix that 'solves' [the least-squares problem] $Ax = b$," i.e., if \bar{x} is said solution, then A^+ is that matrix such that $\bar{x} = A^+b$.

It can be shown that if $Q_1 \Sigma Q_2^T = A$ is the singular value decomposition of A , then $A^+ = Q_2 \Sigma^+ Q_1^T$, where $Q_{1,2}$ are orthogonal matrices, Σ is a diagonal matrix consisting of A 's so-called singular values, (followed, typically, by zeros), and then Σ^+ is simply the diagonal matrix consisting of the reciprocals of A 's singular values (again, followed by zeros). [1]

References

.. [1] G. Strang, *Linear Algebra and Its Applications*, 2nd Ed., Orlando, FL, Academic Press, Inc., 1980, pp. 139-142.

Examples

The following example checks that `a * a+ * a == a` and `a+ * a * a+ == a+`:

```
>>> a = np.random.randn(9, 6)
>>> B = np.linalg.pinv(a)
>>> np.allclose(a, np.dot(a, np.dot(B, a)))
```

```
True
>>> np.allclose(B, np.dot(B, np.dot(a, B)))
True
```

```
qr(a, mode='reduced')
```

Compute the qr factorization of a matrix.

Factor the matrix `a` as qr^* , where `q` is orthonormal and `r` is upper-triangular.

Parameters

`a` : array_like, shape (M, N)

Matrix to be factored.

`mode` : {'reduced', 'complete', 'r', 'raw'}, optional

If $K = \min(M, N)$, then

- * 'reduced' : returns q, r with dimensions (M, K), (K, N) (default)
- * 'complete' : returns q, r with dimensions (M, M), (M, N)
- * 'r' : returns r only with dimensions (K, N)
- * 'raw' : returns h, tau with dimensions (N, M), (K,)

The options 'reduced', 'complete', and 'raw' are new in numpy 1.8, see the notes for more information. The default is 'reduced', and to maintain backward compatibility with earlier versions of numpy both it and the old default 'full' can be omitted. Note that array h returned in 'raw' mode is transposed for calling Fortran. The 'economic' mode is deprecated. The modes 'full' and 'economic' may be passed using only the first letter for backwards compatibility, but all others must be spelled out. See the Notes for more explanation.

Returns

`q` : ndarray of float or complex, optional

A matrix with orthonormal columns. When `mode = 'complete'` the result is an orthogonal/unitary matrix depending on whether or not `a` is real/complex. The determinant may be either ± 1 in that case.

`r` : ndarray of float or complex, optional

The upper-triangular matrix.

`(h, tau)` : ndarrays of np.double or np.cdouble, optional

The array h contains the Householder reflectors that generate q along with r. The tau array contains scaling factors for the

reflectors. In the deprecated 'economic' mode only h is returned.

Raises

LinAlgError

If factoring fails.

Notes

This is an interface to the LAPACK routines ``dgeqrf``, ``zgeqrf``, ``dorgqr``, and ``zungqr``.

For more information on the qr factorization, see for example:

https://en.wikipedia.org/wiki/QR_factorization

Subclasses of `ndarray` are preserved except for the 'raw' mode. So if `a` is of type `matrix`, all the return values will be matrices too.

New 'reduced', 'complete', and 'raw' options for mode were added in NumPy 1.8.0 and the old option 'full' was made an alias of 'reduced'. In addition the options 'full' and 'economic' were deprecated. Because 'full' was the previous default and 'reduced' is the new default, backward compatibility can be maintained by letting `mode` default. The 'raw' option was added so that LAPACK routines that can multiply arrays by q using the Householder reflectors can be used. Note that in this case the returned arrays are of type np.double or np.cdouble and the h array is transposed to be FORTRAN compatible. No routines using the 'raw' return are currently exposed by numpy, but some are available in lapack_lite and just await the necessary work.

Examples

```
>>> a = np.random.randn(9, 6)
>>> q, r = np.linalg.qr(a)
>>> np.allclose(a, np.dot(q, r)) # a does equal qr
True
>>> r2 = np.linalg.qr(a, mode='r')
>>> np.allclose(r, r2) # mode='r' returns the same r as mode='full'
True
```

Example illustrating a common use of `qr`: solving of least squares problems

What are the least-squares-best `m` and `y0` in ``y = y0 + mx`` for the following data: {(0,1), (1,0), (1,2), (2,1)}. (Graph the points

and you'll see that it should be $y_0 = 0$, $m = 1$.) The answer is provided by solving the over-determined matrix equation $Ax = b$, where::

```
A = array([[0, 1], [1, 1], [1, 1], [2, 1]])
x = array([[y0], [m]])
b = array([[1], [0], [2], [1]])
```

If $A = qr$ such that q is orthonormal (which is always possible via Gram-Schmidt), then $x = \text{inv}(r) * (q.T) * b$. (In numpy practice, however, we simply use `lstsq`.)

```
>>> A = np.array([[0, 1], [1, 1], [1, 1], [2, 1]])
>>> A
array([[0, 1],
       [1, 1],
       [1, 1],
       [2, 1]])
>>> b = np.array([1, 0, 2, 1])
>>> q, r = np.linalg.qr(A)
>>> p = np.dot(q.T, b)
>>> np.dot(np.linalg.inv(r), p)
array([ 1.1e-16,  1.0e+00])
```

`slogdet(a)`

Compute the sign and (natural) logarithm of the determinant of an array.

If an array has a very small or very large determinant, then a call to `det` may overflow or underflow. This routine is more robust against such issues, because it computes the logarithm of the determinant rather than the determinant itself.

Parameters

`a` : (...) array_like

Input array, has to be a square 2-D array.

Returns

`sign` : (...) array_like

A number representing the sign of the determinant. For a real matrix, this is 1, 0, or -1. For a complex matrix, this is a complex number with absolute value 1 (i.e., it is on the unit circle), or else 0.

`logdet` : (...) array_like

The natural log of the absolute value of the determinant.

If the determinant is zero, then ``sign`` will be 0 and ``logdet`` will be `-Inf`. In all cases, the determinant is equal to ``sign * np.exp(logdet)``.

See Also

`det`

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

.. versionadded:: 1.6.0

The determinant is computed via LU factorization using the LAPACK routine ``z/dgetrf``.

Examples

The determinant of a 2-D array ``[[a, b], [c, d]]`` is ``ad - bc``:

```
>>> a = np.array([[1, 2], [3, 4]])
>>> (sign, logdet) = np.linalg.slogdet(a)
>>> (sign, logdet)
(-1, 0.69314718055994529) # may vary
>>> sign * np.exp(logdet)
-2.0
```

Computing log-determinants for a stack of matrices:

```
>>> a = np.array([ [[1, 2], [3, 4]], [[1, 2], [2, 1]], [[1, 3], [3, 1]] ])
>>> a.shape
(3, 2, 2)
>>> sign, logdet = np.linalg.slogdet(a)
>>> (sign, logdet)
(array([-1., -1., -1.]), array([ 0.69314718,  1.09861229,  2.07944154]))
>>> sign * np.exp(logdet)
array([-2., -3., -8.])
```

This routine succeeds where ordinary ``det`` does not:

```
>>> np.linalg.det(np.eye(500) * 0.1)
0.0
>>> np.linalg.slogdet(np.eye(500) * 0.1)
(1, -1151.2925464970228)
```

`solve(a, b)`

Solve a linear matrix equation, or system of linear scalar equations.

Computes the "exact" solution, ``x``, of the well-determined, i.e., full rank, linear matrix equation ``ax = b``.

Parameters

`a` : (... , M, M) array_like
Coefficient matrix.
`b` : {... , M,), (... , M, K)}, array_like
Ordinate or "dependent variable" values.

Returns

`x` : {... , M,), (... , M, K)} ndarray
Solution to the system `a x = b`. Returned shape is identical to ``b``.

Raises

`LinAlgError`
If ``a`` is singular or not square.

Notes

.. versionadded:: 1.8.0

Broadcasting rules apply, see the ``numpy.linalg`` documentation for details.

The solutions are computed using LAPACK routine ```_gesv```.

``a`` must be square and of full-rank, i.e., all rows (or, equivalently, columns) must be linearly independent; if either is not true, use ``lstsq`` for the least-squares best "solution" of the system/equation.

References

```
.. [1] G. Strang, *Linear Algebra and Its Applications*, 2nd Ed., Orlando,
    FL, Academic Press, Inc., 1980, pg. 22.
```

Examples

Solve the system of equations $3 * x_0 + x_1 = 9$ and $x_0 + 2 * x_1 = 8$:

```
>>> a = np.array([[3,1], [1,2]])
>>> b = np.array([9,8])
>>> x = np.linalg.solve(a, b)
>>> x
array([2., 3.] )
```

Check that the solution is correct:

```
>>> np.allclose(np.dot(a, x), b)
True
```

```
svd(a, full_matrices=True, compute_uv=True, hermitian=False)
Singular Value Decomposition.
```

When `a` is a 2D array, it is factorized as `u @ np.diag(s) @ vh`
= `(u * s) @ vh`, where `u` and `vh` are 2D unitary arrays and `s` is a 1D
array of `a`'s singular values. When `a` is higher-dimensional, SVD is
applied in stacked mode as explained below.

Parameters

`a` : (... , M, N) array_like
A real or complex array with `a.ndim >= 2`.
`full_matrices` : bool, optional
If True (default), `u` and `vh` have the shapes `((... , M, M))` and
`((... , N, N))`, respectively. Otherwise, the shapes are
`((... , M, K))` and `((... , K, N))`, respectively, where
`K = min(M, N)`.
`compute_uv` : bool, optional
Whether or not to compute `u` and `vh` in addition to `s`. True
by default.
`hermitian` : bool, optional
If True, `a` is assumed to be Hermitian (symmetric if real-valued),
enabling a more efficient method for finding singular values.
Defaults to False.

.. versionadded:: 1.17.0

Returns

`u` : { (... , M, M), (... , M, K) } array
Unitary array(s). The first ```a.ndim - 2``` dimensions have the same size as those of the input `a`. The size of the last two dimensions depends on the value of `full_matrices`. Only returned when `compute_uv` is True.

`s` : (... , K) array
Vector(s) with the singular values, within each vector sorted in descending order. The first ```a.ndim - 2``` dimensions have the same size as those of the input `a`.

`vh` : { (... , N, N), (... , K, N) } array
Unitary array(s). The first ```a.ndim - 2``` dimensions have the same size as those of the input `a`. The size of the last two dimensions depends on the value of `full_matrices`. Only returned when `compute_uv` is True.

Raises

LinAlgError

If SVD computation does not converge.

Notes

.. versionchanged:: 1.8.0

Broadcasting rules apply, see the `numpy.linalg` documentation for details.

The decomposition is performed using LAPACK routine ```_gesdd```.

SVD is usually described for the factorization of a 2D matrix A . The higher-dimensional case will be discussed below. In the 2D case, SVD is written as $A = U S V^H$, where $A = a$, $U = u$, $S = \text{np.diag}(s)$ and $V^H = vh$. The 1D array `s` contains the singular values of `a` and `u` and `vh` are unitary. The rows of `vh` are the eigenvectors of $A^H A$ and the columns of `u` are the eigenvectors of $A A^H$. In both cases the corresponding (possibly non-zero) eigenvalues are given by `s**2`.

If `a` has more than two dimensions, then broadcasting rules apply, as explained in :ref:`routines.linalg-broadcasting`. This means that SVD is working in "stacked" mode: it iterates over all indices of the first ```a.ndim - 2``` dimensions and for each combination SVD is applied to the last two indices. The matrix `a` can be reconstructed from the

decomposition with either `((u * s[..., None, :]) @ vh)` or `((u @ (s[..., None] * vh))`. (The `@` operator can be replaced by the function `np.matmul` for python versions below 3.5.)

If `a` is a `matrix` object (as opposed to an `ndarray`), then so are all the return values.

Examples

```
>>> a = np.random.randn(9, 6) + 1j*np.random.randn(9, 6)
>>> b = np.random.randn(2, 7, 8, 3) + 1j*np.random.randn(2, 7, 8, 3)
```

Reconstruction based on full SVD, 2D case:

```
>>> u, s, vh = np.linalg.svd(a, full_matrices=True)
>>> u.shape, s.shape, vh.shape
((9, 9), (6,), (6, 6))
>>> np.allclose(a, np.dot(u[:, :6] * s, vh))
True
>>> smat = np.zeros((9, 6), dtype=complex)
>>> smat[:6, :6] = np.diag(s)
>>> np.allclose(a, np.dot(u, np.dot(smat, vh)))
True
```

Reconstruction based on reduced SVD, 2D case:

```
>>> u, s, vh = np.linalg.svd(a, full_matrices=False)
>>> u.shape, s.shape, vh.shape
((9, 6), (6,), (6, 6))
>>> np.allclose(a, np.dot(u * s, vh))
True
>>> smat = np.diag(s)
>>> np.allclose(a, np.dot(u, np.dot(smat, vh)))
True
```

Reconstruction based on full SVD, 4D case:

```
>>> u, s, vh = np.linalg.svd(b, full_matrices=True)
>>> u.shape, s.shape, vh.shape
((2, 7, 8, 8), (2, 7, 3), (2, 7, 3, 3))
>>> np.allclose(b, np.matmul(u[..., :3] * s[..., None, :], vh))
True
>>> np.allclose(b, np.matmul(u[..., :3], s[..., None] * vh))
True
```

Reconstruction based on reduced SVD, 4D case:

```
>>> u, s, vh = np.linalg.svd(b, full_matrices=False)
>>> u.shape, s.shape, vh.shape
((2, 7, 8, 3), (2, 7, 3), (2, 7, 3, 3))
>>> np.allclose(b, np.matmul(u * s[..., None, :], vh))
True
>>> np.allclose(b, np.matmul(u, s[..., None] * vh))
True
```

`tensorinv(a, ind=2)`

Compute the 'inverse' of an N-dimensional array.

The result is an inverse for ``a`` relative to the `tensor` operation ``tensor(a, b, ind)``, i. e., up to floating-point accuracy, ``tensor(tensorinv(a), a, ind)`` is the "identity" tensor for the `tensor` operation.

Parameters

`a` : array_like

Tensor to 'invert'. Its shape must be 'square', i. e., ``prod(a.shape[:ind]) == prod(a.shape[ind:])``.

`ind` : int, optional

Number of first indices that are involved in the inverse sum. Must be a positive integer, default is 2.

Returns

`b` : ndarray

``a``'s `tensor` inverse, shape ``a.shape[ind:] + a.shape[:ind]``.

Raises

`LinAlgError`

If ``a`` is singular or not 'square' (in the above sense).

See Also

`numpy.tensordot`, `tensorsolve`

Examples

```
>>> a = np.eye(4*6)
>>> a.shape = (4, 6, 8, 3)
```

```

>>> ainv = np.linalg.tensorinv(a, ind=2)
>>> ainv.shape
(8, 3, 4, 6)
>>> b = np.random.randn(4, 6)
>>> np.allclose(np.tensordot(ainv, b), np.linalg.tensorsolve(a, b))
True

>>> a = np.eye(4*6)
>>> a.shape = (24, 8, 3)
>>> ainv = np.linalg.tensorinv(a, ind=1)
>>> ainv.shape
(8, 3, 24)
>>> b = np.random.randn(24)
>>> np.allclose(np.tensordot(ainv, b, 1), np.linalg.tensorsolve(a, b))
True

```

`tensorsolve(a, b, axes=None)`
 Solve the tensor equation ```a x = b``` for `x`.

It is assumed that all indices of ```x``` are summed over in the product, together with the rightmost indices of ```a```, as is done in, for example, ```tensordot(a, x, axes=b.ndim)```.

Parameters

`a` : array_like
 Coefficient tensor, of shape ```b.shape + Q```. ```Q```, a tuple, equals the shape of that sub-tensor of ```a``` consisting of the appropriate number of its rightmost indices, and must be such that ```prod(Q) == prod(b.shape)``` (in which sense ```a``` is said to be 'square').

`b` : array_like
 Right-hand tensor, which can be of any shape.

`axes` : tuple of ints, optional
 Axes in ```a``` to reorder to the right, before inversion.
 If None (default), no reordering is done.

Returns

`x` : ndarray, shape `Q`

Raises

`LinAlgError`
 If ```a``` is singular or not 'square' (in the above sense).

See Also

numpy.tensordot, tensorinv, numpy.einsum

Examples

```
>>> a = np.eye(2*3*4)
>>> a.shape = (2*3, 4, 2, 3, 4)
>>> b = np.random.randn(2*3, 4)
>>> x = np.linalg.tensorsolve(a, b)
>>> x.shape
(2, 3, 4)
>>> np.allclose(np.tensordot(a, x, axes=3), b)
True
```

DATA

```
absolute_import = _Feature((2, 5, 0, 'alpha', 1), (3, 0, 0, 'alpha', 0...
division = _Feature((2, 2, 0, 'alpha', 2), (3, 0, 0, 'alpha', 0), 8192...
print_function = _Feature((2, 6, 0, 'alpha', 2), (3, 0, 0, 'alpha', 0)...
test = <numpy._pytesttester.PytestTester object>
```

FILE

/Users/b0a00c8/Documents/work3/venv_new/lib/python3.7/site-packages/numpy/linalg/__init__.py

```
In [74]: A = np.array([[6, 1, 1],
                      [4, -2, 5],
                      [2, 8, 7]])
```

```
In [75]: A
```

```
Out[75]: array([[ 6,  1,  1],
                [ 4, -2,  5],
                [ 2,  8,  7]])
```

Rank of a matrix

```
In [76]: np.linalg.matrix_rank(A)
```

```
Out[76]: 3
```


Trace of matrix A

```
In [77]: np.trace(A)
```

```
Out[77]: 11
```

Determinant of a matrix

```
In [78]: np.linalg.det(A)
```

```
Out[78]: -306.0
```

Inverse of matrix A

```
In [87]: A
```

```
Out[87]: array([[ 6,  1,  1],
                [ 4, -2,  5],
                [ 2,  8,  7]])
```

```
In [79]: np.linalg.inv(A)
```

```
Out[79]: array([[ 0.17647059, -0.00326797, -0.02287582],
                [ 0.05882353, -0.13071895,  0.08496732],
                [-0.11764706,  0.1503268 ,  0.05228758]])
```

```
In [84]: B = np.linalg.inv(A)
```

```
In [85]: np.matmul(A,B) #actual matrix multiplication
```

```
Out[85]: array([[ 1.00000000e+00,  0.00000000e+00,  2.77555756e-17],
                [-1.38777878e-17,  1.00000000e+00,  1.38777878e-17],
                [-4.16333634e-17,  1.38777878e-16,  1.00000000e+00]])
```

```
In [86]: A * B
```

```
Out[86]: array([[ 1.05882353, -0.00326797, -0.02287582],
                [ 0.23529412,  0.26143791,  0.4248366 ],
                [-0.23529412,  1.20261438,  0.36601307]])
```

Matrix A raised to power 3

```
In [88]: np.linalg.matrix_power(A,3) # matrix multiplication A A A
```

```
Out[88]: array([[336, 162, 228],  
               [406, 162, 469],  
               [698, 702, 905]])
```

```
In [ ]:
```